

# KURSE: Knowledge Utilization, Retrieval and Summarization Engine

with Application to BeeKurse E-Commerce Platform

Varshith Kada      Vishnu Teja      Himesh Anant      Siripuru Abhiram  
Abhinav Goyal      Bhuvan Bhedika      Kunjan Manoj

## Abstract

KURSE is a neuro-symbolic inference engine combining neural networks with symbolic reasoning through a tri-store architecture (SQL, Vector DB, Knowledge Graph). We present the KG ingestion pipeline with N-ary event representation and demonstrate two applications: *Alibi Breaker* for constraint satisfaction and *BeeKurse* for conversational e-commerce. BeeKurse features a 5-type query classifier, multi-strategy search orchestration (HQ/RQ/SQ paths), and intelligent scoring with property weights (0.5-2.0 scale).

## 1 Introduction

Modern applications require systems that can both store and reason over unstructured data. KURSE addresses this through a **neuro-symbolic approach**:

- **Neural:** LLMs for entity extraction and semantic parsing
- **Symbolic:** Knowledge graphs with ontologies for structured reasoning

### 1.1 Tri-Store Architecture

KURSE integrates three storage paradigms:

1. **SQL:** Structured metadata with ACID guarantees
2. **Vector DB (Qdrant):** Semantic similarity via 1024-dim NVCLIP embeddings
3. **Knowledge Graph (Memgraph):** Entity relationships with Cypher queries

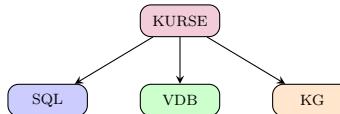


Figure 1: Tri-Store Architecture

## 2 KURSE Core Pipeline

### 2.1 Input Processing

Documents undergo OCR (olmOCR/Qwen 7B), then chunking with parameters:

Parameter	Value
CHUNK_SIZE	1000 chars
CHUNK_OVERLAP	200 chars
NVCLIP_DIM	1024

## 2.2 KG Ingestion Pipeline

The pipeline operates in two phases:

**Phase 1 - N-ary Event Identification:** Identifies complex events involving 3+ participants or critical metadata (time, location). Events are stored with embeddings for semantic retrieval.

**Phase 2 - Relation Extraction:** Extracts triplets (subject-predicate-object), verifies against source text, and standardizes to ontology using VDB similarity:

Threshold	Value
TYPE_THRESHOLD	0.25
RELATION_THRESHOLD	0.20

A streaming bridge connects phases via async queues, enabling parallel processing before Memgraph storage.

## 2.3 Ontology & Reasoning

KURSE supports formal axioms: **transitivity** ( $R(x, y) \wedge R(y, z) \Rightarrow R(x, z)$ ), **symmetry**, and **domain/range constraints** for validation. See Appendix D for details.

## 2.4 Applications

KURSE is at its core an intelligent data storage hence it's ubiquitously applicable. Let's see some applications to see the potential of what we can do with the KURSE Engine.

**Alibi Breaker:** Constraint satisfaction to detect contradictions in statements with verified facts or evidence.

**Risk Contagion:** Identify Non-Obvious risk to the CCC, cash flow operations or stock price with global news integration to identify indirect impacts and system risk.

**BeeKurse:** A Conversational E commerce Agent specializing in product recommendation

## 3 BeeKurse: E-Commerce Application

BeeKurse adapts KURSE for conversational product search, dropping strict ontology for flexible VDB-based relations.

### 3.1 Query Classification

The **Strontium** parser (269-line system prompt) classifies into 5 types:

### 3.2 SEARCH Query Model

Parsed into **ProductRequest** with:

- **Properties:** (name, value, weight) where weight  $\in [0.5, 2.0]$

Type	Purpose
SEARCH	Product search with properties/literals
DETAIL	Info about specific product
CHAT	General conversation
CART_ACTION	Add/remove from cart
CART_VIEW	View cart contents

- **Literals:** Numeric constraints with buffers
- **is\_hq:** Hurry Query flag for fast path
- **sort\_literal:** For superlatives (“cheapest”)

### 3.3 Search Orchestration

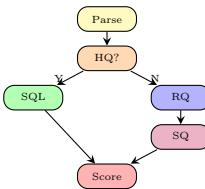


Figure 2: Search Pipeline

**HQ Fast Path:** Direct SQL for urgent/generic queries.

**Property Search (RQ):** Dual VDB search (main + property collections) with semantic matching.

**Connected Search (SQ):** KG traversal via ALSO\_BOUGHT, SIMILAR\_TO, COMPLEMENTS relations.

### 3.4 Scoring System

$$\text{score} = \sum_i (sim_i \times w_i) + B_{conn} + B_{subcat}$$

Where  $B_{conn} = 0.5$  (connected bonus),  $B_{subcat} = 0.4$  (subcategory bonus).

For superlatives: 70% literal ranking + 30% semantic ranking.

## 4 Conclusion

KURSE demonstrates effective neuro-symbolic integration through tri-store architecture, N-ary event representation, and ontology-driven reasoning. BeeKurse showcases adaptability with sophisticated query classification, multi-strategy search, and intelligent scoring.

## References

- [1] Aidan Hogan, Eva Blomqvist, Michael Cochez, Claudia d’Amato, Gerard de Melo, Claudio Gutiérrez, Sabrina Kirrane, Jose Emilio Labra-Gayo, Roberto Navigli, Sebastian Neu-maier, A. Ngonga Moses Ngomo, Axel Polleres, Shanmugalingam Rashid, Anisa Rula, Lukas Schmelzeisen, Juan F. Sequeda, Steffen Staab, and Antoine Zimmermann. Knowledge graphs. *ACM Computing Surveys*, 54(4):1–37, 2021.

## Appendices

### A Input Processing Details

#### A.1 OCR with olmOCR

The OCR module uses **olmOCR**, a vision-language model based on Qwen 7B, fine-tuned for:

- Dense text recognition in documents
- Table structure understanding
- Handwritten text interpretation
- Multi-language support

#### A.2 Document ID Generation

Each document receives a unique identifier combining content hash and timestamp:

```
import hashlib
from datetime import datetime

def generate_document_id(content: bytes) -> str:
    content_hash = hashlib.md5(content).hexdigest()[:8]
    timestamp = datetime.now().strftime("%Y%m%d%H%M%S")
    return f"doc_{content_hash}_{timestamp}"
```

#### A.3 Chunking Algorithm

The chunking algorithm respects sentence boundaries:

```
def chunk_document(text, chunk_size=1000, overlap=200):
    chunks = []
    start = 0
    while start < len(text):
        end = min(start + chunk_size, len(text))
        if end < len(text):
            # Find sentence boundary
            boundary = find_sentence_boundary(text[start:end])
            if boundary > chunk_size * 0.7:
                end = start + boundary
        chunks.append(text[start:end])
        start = end - overlap
    return chunks
```

## B Vector Database Schema

### B.1 NVCLIP Embeddings

KURSE uses NVIDIA's NVCLIP model:

- **Model:** nvidia/nvclip
- **Dimension:** 1024
- **Normalization:** L2-normalized for cosine similarity

### B.2 Qdrant Point Structure

```
{
  "id": "chunk_abcd123_001",
  "vector": [0.123, -0.456, ...], # 1024 dimensions
  "payload": {
    "document_id": "doc_a1b2c3d4_20240101120000",
    "chunk_id": "chunk_001",
    "chunk_index": 0,
    "page_no": 1,
    "text": "Original chunk text...",
    "created_at": "2024-01-01T12:00:00Z"
  }
}
```

### B.3 Collection Configuration

```
from qdrant_client import QdrantClient
from qdrant_client.models import VectorParams, Distance

client.create_collection(
    collection_name="document_chunks",
    vectors_config=VectorParams(
        size=1024,
        distance=Distance.COSINE
    )
)
```

## C Knowledge Graph Fundamentals

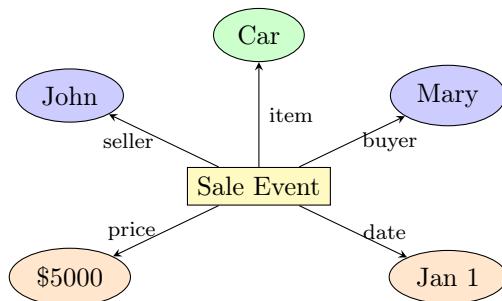
### C.1 Node Types

The KURSE knowledge graph has three node categories:

1. **Entities**: Real-world objects (Person, Organization, Location, Event)
2. **Properties**: Attributes (color, size, material, date)
3. **Literals**: Concrete values (“red”, 42, “2024-01-01”)

### C.2 N-ary Relations

Traditional binary triplets are insufficient for complex events. N-ary relations create an **Event Node** connecting all participants:



### C.3 Cypher Schema

```
// Create Event Node
CREATE (e:Event {
  id: "event_001",
```

```

        type: "Sale",
        source_chunk: "chunk_abc123_001",
        confidence: 0.95
    })

// Connect Participants
MATCH (e:Event {id: "event_001"})
MATCH (seller:Person {name: "John"})
MATCH (buyer:Person {name: "Mary"})
CREATE (e)-[:HAS_SELLER]->(seller)
CREATE (e)-[:HAS_BUYER]->(buyer)

```

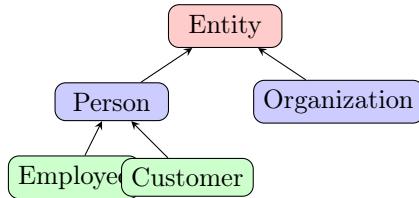
## D Ontology & Reasoning

### D.1 What is an Ontology?

An **ontology** is a formal, explicit specification of a shared conceptualization:

- **Formal:** Machine-readable with precise semantics
- **Explicit:** Clearly defined concepts and constraints
- **Shared:** Common understanding across components

### D.2 Class Hierarchies (subClassOf)



### D.3 Ontology Axioms

**Transitivity:**

$$\forall x, y, z : R(x, y) \wedge R(y, z) \Rightarrow R(x, z)$$

Example: LOCATED\_IN is transitive.

**Symmetry:**

$$\forall x, y : R(x, y) \Rightarrow R(y, x)$$

Example: SIBLING\_OF is symmetric.

**Equivalence:**

$$A \equiv B \Leftrightarrow (A \sqsubseteq B) \wedge (B \sqsubseteq A)$$

### D.4 First-Order Logic Mappings

### D.5 Domain and Range Constraints

```

def validate_relation(subject, predicate, object, ontology):
    """Validate relation against domain/range constraints."""
    domain = ontology.get_domain(predicate)
    range_ = ontology.get_range(predicate)

    if not is_instance_of(subject, domain):
        raise ValidationError(f"Domain violation")
    if not is_instance_of(object, range_):

```

Ontology Construct	FOL Representation
subClassOf(A, B)	$\forall x : A(x) \Rightarrow B(x)$
domain(R, C)	$\forall x, y : R(x, y) \Rightarrow C(x)$
range(R, C)	$\forall x, y : R(x, y) \Rightarrow C(y)$
transitive(R)	$\forall x, y, z : R(x, y) \wedge R(y, z) \Rightarrow R(x, z)$
symmetric(R)	$\forall x, y : R(x, y) \Rightarrow R(y, x)$

```
    raise ValidationError(f"Range violation")
    return True
```

## E KG Ingestion Pipeline (Full)

### E.1 Phase 1: N-ary Event Identification

Event Criteria:

1. **Three or more participants:** Multiple entities in single action
2. **Critical metadata:** Temporal, spatial, or quantitative data
3. **Semantic unity:** Participants connected by coherent action

Event Data Structure:

```
@dataclass
class NaryEvent:
    event_id: str
    event_type: str
    participants: List[Participant]
    metadata: Dict[str, Any]
    source_chunk_id: str
    source_document_id: str
    confidence: float
    embedding: List[float] # 1024-dim

@dataclass
class Participant:
    entity_name: str
    entity_type: str
    role: str # seller, buyer, witness, etc.
```

Manifest Structure:

```
{
  "document_id": "doc_abc123",
  "chunk_id": "chunk_001",
  "page_no": 1,
  "events_identified": [
    {"event_id": "evt_001", "event_type": "Transaction"}
  ],
  "processing_status": "phase1_complete"
}
```

### E.2 Phase 2: Relation Extraction

RawTriplet Structure:

```
class RawTriplet(BaseModel):
    subject: str # Entity name
    subject_type: str # Raw type from LLM
```

```

predicate: str          # Relation name
object: str             # Entity or literal
object_type: str         # Raw type from LLM
confidence: float        # Extraction confidence
source_span: Optional[str]

```

### Verification Against Source:

```

def verify_triplet(triplet: RawTriplet, source_text: str) -> bool:
    """Verify triplet is supported by source text."""
    verification_prompt = f"""
Source text: {source_text}
Claimed fact: {triplet.subject}{triplet.predicate}{triplet.object}
Is this fact directly supported by the source text?
"""
    response = llm.generate(verification_prompt)
    return response.startswith("YES")

```

### Standardization to Ontology:

```

TYPE_THRESHOLD = 0.25
RELATION_THRESHOLD = 0.20

def standardize_type(raw_type: str, types_vdb) -> str:
    """Match raw type to ontology using VDB."""
    embedding = encoder.encode(raw_type)
    results = types_vdb.search(
        collection_name="entity_types",
        query_vector=embedding,
        limit=1
    )
    if results[0].score >= TYPE_THRESHOLD:
        return results[0].payload["type_name"]
    return "Unknown"

```

## E.3 Streaming Bridge Architecture

```

class StreamingBridge:
    def __init__(self, num_workers: int = 4):
        self.queue = Queue()
        self.num_workers = num_workers

    async def phase1_producer(self, chunks):
        for chunk in chunks:
            events = await identify_nary_events(chunk)
            manifest = create_manifest(chunk, events)
            await self.queue.put((chunk, manifest))
        for _ in range(self.num_workers):
            await self.queue.put(None)

    async def phase2_worker(self, worker_id):
        while True:
            item = await self.queue.get()
            if item is None: break
            chunk, manifest = item
            triplets = await extract_relations(chunk, manifest)
            await store_to_memgraph(triplets)

```

## E.4 Memgraph Storage

```

// Store Entity (idempotent)
MERGE (e:Entity {name: $name})
ON CREATE SET
    e.type = $type,
    e.created_at = timestamp(),
    e.source_documents = [$document_id]
ON MATCH SET
    e.source_documents = e.source_documents + $document_id

```

```

// Store Relation
MATCH (s:Entity {name: $subject})
MATCH (o:Entity {name: $object})
MERGE (s)-[r:$predicate]->(o)
ON CREATE SET
    r.source_chunk = $chunk_id,
    r.confidence = $confidence

```

## F Verification System

### F.1 Chunk Verification Pipeline

```

async def verify_chunk_extractions(chunk, triplets):
    """Verify extracted triplets against source chunk."""
    verified = []
    for triplet in triplets:
        prompt = f"""
Source;text:{chunk}
Verify this extracted fact:
Subject:{triplet.subject}(type:{triplet.subject_type})
Relation:{triplet.predicate}
Object:{triplet.object}(type:{triplet.object_type})

Is this fact:
1. Explicitly stated in the text?
2. Reasonably inferable from the text?
3. Not supported by the text?
"""

        response = await llm.generate(prompt)
        if response.startswith("1") or response.startswith("2"):
            verified.append(triplet)
    return verified

```

### F.2 Ontology Feasibility Checking

```

def check_ontology_feasibility(triplet, ontology):
    """Check if triplet is feasible under ontology."""
    if triplet.predicate in ontology.relations:
        expected_domain = ontology.get_domain(triplet.predicate)
        if not ontology.is_subclass(triplet.subject_type, expected_domain):
            return False, f"Domain violation"
        expected_range = ontology.get_range(triplet.predicate)
        if not ontology.is_subclass(triplet.object_type, expected_range):
            return False, f"Range violation"
    return True, "Valid"

```

## G Query Classification Details

### G.1 Query Type Output Formats

#### SEARCH Query Output:

```
{
    "query_type": "SEARCH",
    "product_request": {
        "product_category": "dress",
        "properties": [
            {"name": "color", "value": "red", "weight": 1.5}
        ],
        "literals": [
            {"name": "price", "value": 50, "operator": "<=", "buffer": 7.5}
        ],
    }
}
```

```

        "is_hq": false,
        "sort_literal": null
    }
}

```

### DETAIL Query Output:

```
{
    "query_type": "DETAIL",
    "product_id": "123",
    "question": "What material is this made of?"
}
```

### CART\_ACTION Output:

```
{
    "query_type": "CART_ACTION",
    "action": "add",
    "product_id": "123",
    "quantity": 2
}
```

## G.2 ProductRequest Model

```

class Property(BaseModel):
    name: str
    value: str
    weight: float = 1.0 # 0.5 to 2.0 scale

class Literal(BaseModel):
    name: str
    value: float
    operator: str # =, <, <=, >, >=, between
    buffer: float = 0.0

class ProductRequest(BaseModel):
    product_category: str
    product_subcategory: Optional[str] = None
    properties: List[Property] = []
    literals: List[Literal] = []
    is_hq: bool = False
    sort_literal: Optional[str] = None

```

## G.3 Property Weight Scale

Weight	Level	Indicators
2.0	Critical	“must be”, “definitely”, “nothing else”
1.5	Important	Direct mention with emphasis
1.0	Standard	Simple mention
0.5	Nice-to-have	“preferably”, “if possible”

## H Search Orchestration Details

### H.1 HQ Fast Path

**What:** Direct SQL lookup bypassing semantic search.

**When:**

- User indicates urgency (“show me any”, “quick”)

- Generic queries without specific requirements
- Re-queries after failed detailed search

**Why:** Sub-100ms response, no embedding computation needed.

```
async def hq_fast_path(request):
    query = """
SELECT * FROM products
WHERE category=:category AND status='active'
ORDER BY popularity DESC LIMIT 20
"""
    return await sql_client.fetch_all(
        query, {"category": request.product_category}
    )
```

## H.2 Property Search (RQ)

**What:** Dual-path semantic search using Main VDB and Property VDB.

**When:** User specifies properties (color, material, style).

**Why:** Handles synonyms (“crimson” → “red”), semantic similarity.

```
async def property_search(request):
    results = []
    # Path 1: Main VDB search
    query_embedding = encoder.encode(request.to_search_string())
    main_results = await main_vdb.search(
        collection_name="products",
        query_vector=query_embedding, limit=50
    )
    # Path 2: Property VDB search
    for prop in request.properties:
        prop_embedding = encoder.encode(f"{prop.name}:{prop.value}")
        prop_results = await property_vdb.search(
            collection_name="product_properties",
            query_vector=prop_embedding, limit=30
        )
        for r in prop_results:
            r.score *= prop.weight
        results.extend(prop_results)
    return merge_results(main_results, results)
```

## H.3 Connected Search (SQ)

**What:** Knowledge graph traversal for related products.

**When:** Recommendations, “similar to” queries, complementary products.

**Why:** Discovers non-obvious connections via co-purchase patterns.

```
async def connected_search(seed_products, relation_types=None):
    if relation_types is None:
        relation_types = ["ALSO_BOUGHT", "SIMILAR_TO", "COMPLEMENTS"]
    query = """
MATCH (seed:Product)-[r]-(connected:Product)
WHERE seed.id IN $seed_ids AND type(r) IN $relations
RETURN connected, type(r) as relation, r.strength
ORDER BY r.strength DESC LIMIT 30
"""
    return await memgraph.execute(query, {
        "seed_ids": seed_products, "relations": relation_types
    })
```

## H.4 Subcategory Scoring

```

def apply_subcategory_bonus(products, target_subcategory):
    target_embedding = encoder.encode(target_subcategory)
    for product in products:
        subcat_embedding = encoder.encode(product.subcategory)
        similarity = cosine_similarity(target_embedding, subcat_embedding)
        bonus = similarity * SUBCATEGORY_MAX_BONUS # 0.4
        product.score += bonus
    return products

```

## I Weighting System

### I.1 Property Weight Formula

$$score_{property} = similarity \times weight$$

Where similarity is cosine similarity (0-1) and weight is user importance (0.5-2.0).

### I.2 Score Combination Formula

$$final\_score = \sum_i (property\_scores_i) + connected\_bonus + subcategory\_bonus$$

```

def calculate_final_score(product, property_scores,
                         connected_bonus, subcategory_bonus,
                         literal_results):
    base_score = sum(property_scores.values())
    total_score = base_score + connected_bonus + subcategory_bonus

    for literal_name, passed in literal_results.items():
        if not passed:
            return -1 # Filter out
    return total_score

```

### I.3 Bonus Values

Factor	Value	Description
CONNECTED_BONUS	+0.5	Product found via KG traversal
SUBCATEGORY_MAX_BONUS	+0.4	Max subcategory similarity bonus
EXACT_MATCH_BONUS	+0.25	Exact property value match
LITERAL_FAIL	Filter out	Product doesn't meet constraint

### I.4 Superlative Re-ranking (70/30 Split)

```

def rerank_for_superlative(products, sort_literal, ascending):
    # Sort by literal value
    sorted_by_literal = sorted(
        products, key=lambda p: p.product[sort_literal],
        reverse=not ascending
    )
    literal_ranks = {p.product.id: i for i, p in enumerate(sorted_by_literal)}
    semantic_ranks = {p.product.id: i for i, p in enumerate(products)}

    # Combine: 70% literal, 30% semantic

```

```

for product in products:
    pid = product.product.id
    combined_rank = 0.7 * literal_ranks[pid] + 0.3 * semantic_ranks[pid]
    product.final_rank = combined_rank
return sorted(products, key=lambda p: p.final_rank)

```

## J User Context Enrichment

### J.1 Three-Stage Pipeline

1. **Parse:** Extract user preferences from profile
2. **Filter:** Remove irrelevant preferences for current query
3. **Enrich:** Merge relevant preferences into query with low weight

### J.2 User Context Model

```

class UserContext(BaseModel):
    user_id: str
    gender: Optional[str] = None
    size_preferences: Dict[str, str] = {} # category -> size
    age_group: Optional[str] = None
    style_preferences: List[str] = []
    price_range: Optional[Tuple[float, float]] = None
    brand_preferences: List[str] = []
    recent_searches: List[str] = []
    viewed_products: List[str] = []
    purchased_products: List[str] = []

```

### J.3 Property Merging Strategy

```

def merge_user_context(request, context):
    enriched = request.copy()

    # Add gender if relevant and not specified
    if context.gender and not has_property(request, "gender"):
        if is_gendered_category(request.product_category):
            enriched.properties.append(
                Property(name="gender", value=context.gender, weight=0.5)
            )

    # Add size preference
    if request.product_category in context.size_preferences:
        if not has_property(request, "size"):
            enriched.properties.append(Property(
                name="size",
                value=context.size_preferences[request.product_category],
                weight=0.5
            ))

    # Add style preferences with low weight
    for style in context.style_preferences[:2]:
        if not has_property(request, "style"):
            enriched.properties.append(
                Property(name="style", value=style, weight=0.3)
            )

    return enriched

```