

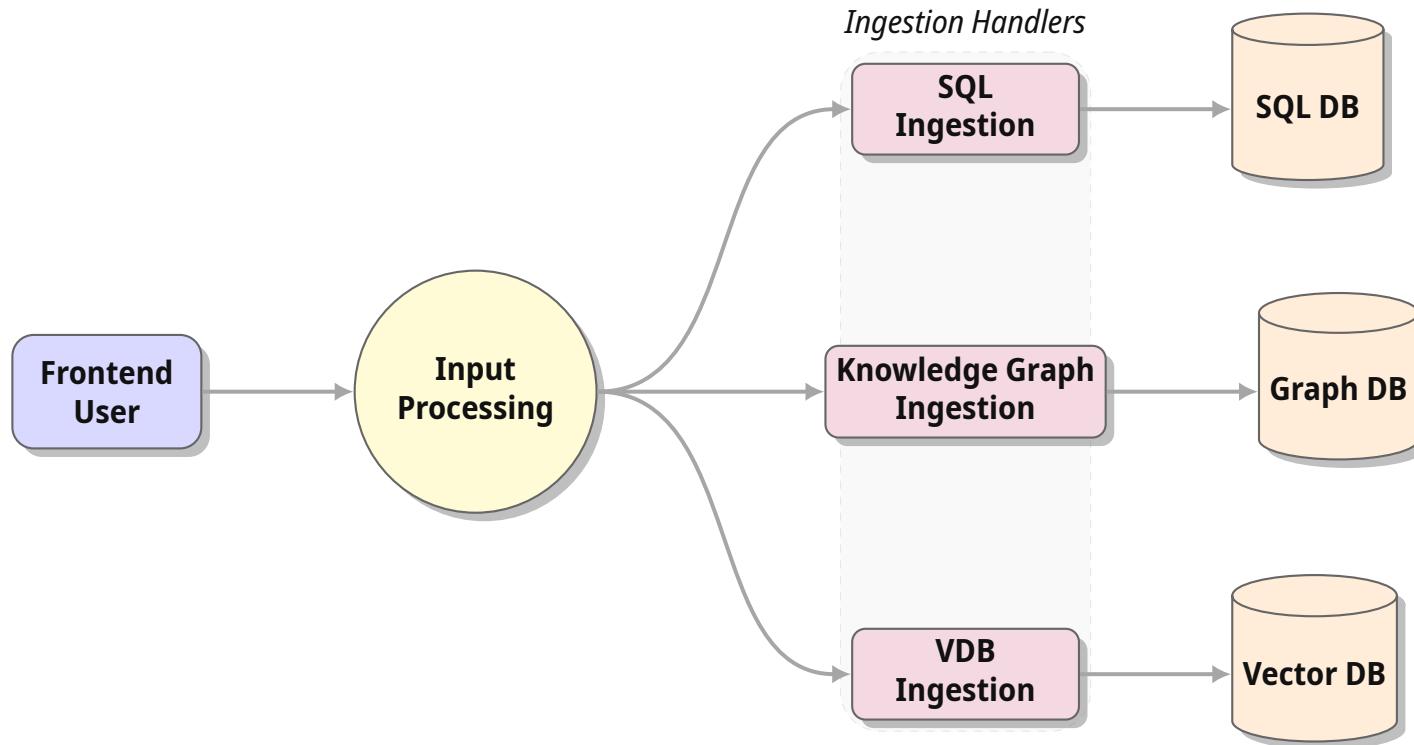
# KURSE

\*built entirely with open source models and infrastructure

# What is KURSE?

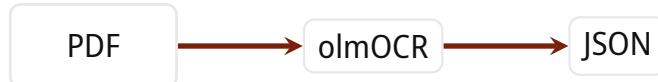
- > Neuro-Symbolic Inference Engine
- > Intelligent tri-store architecture: SQL, Vector DB, and Knowledge Graph
- > Ontology-driven structure enabling hierarchical reasoning for events & relations

Let's see how Kurse works



# Input Processing

- > Takes any unreadable data such as handwritten documents, PDFs and converts them into a readable format
- > **Engine:** Powered by **olmOCR** (based on Qwen 7B)
- > **Goal:** Structured JSON output optimized for LLM readability
- > **Complex Element Handling:**
  - > Tables & Complex Structures → HTML
  - > Standard Text → Markdown



# Input Processing

- > **Metadata Injection:**

- document\_id, page\_no, char\_count

- > **Performance Metrics:**

- > Latency: ~1.2s per page
  - > Supports Parallel Batch Ingestion

*\*Metrics based on local running with a 4090 GPU*

# Vector Database Ingestion

- > **Model Architecture:** Uses **CLIP**
  - > Creates a **shared vector space** for both images and text
  - > Enables seamless cross-modal search and retrieval
- > **Embedding Specification:** High-fidelity **1024** Dimensions
- > **Stored Metadata Payload:**
  - > document\_id, chunk\_id, page\_no
  - > source\_file, text\_preview
  - > table\_no, image\_no

# Knowledge Graph

> Nodes connected by edges (relations)

> **Structure:**

Node(type, props) — Relation(type, origin, version) — Node(type, props)

> **Example:**

(Person: "John") — WORKS\_AT(src: doc\_01, v: 1.2) — (Org: "Acme Corp")

# Node Types

- > **Entities** — concrete things: products, persons, organizations
- > **Properties** — descriptive attributes: material, style, designation
- > **Literals** — terminal values with no relational value: price, date

# N-ary Relations

- > Relations involving **3+ participants** that can't be reduced to binary pairs
- > **Why important:** Complex events need a central node to capture context (time, location, method) that belongs to the event itself, not individual actors

## Example:

(Person: "Marcus") -[COMMITTED]- (Crime: "Heist\_NYC\_2024")  
    -[LOCATION]- (Place: "Manhattan")  
        -[WITH]- (Person: "Sofia")  
    -[ON\_DATE]- (Date: "2024-03-15")

# What is an Ontology?

- A predefined set of **things** with properties & relations

**Entity** (Class)

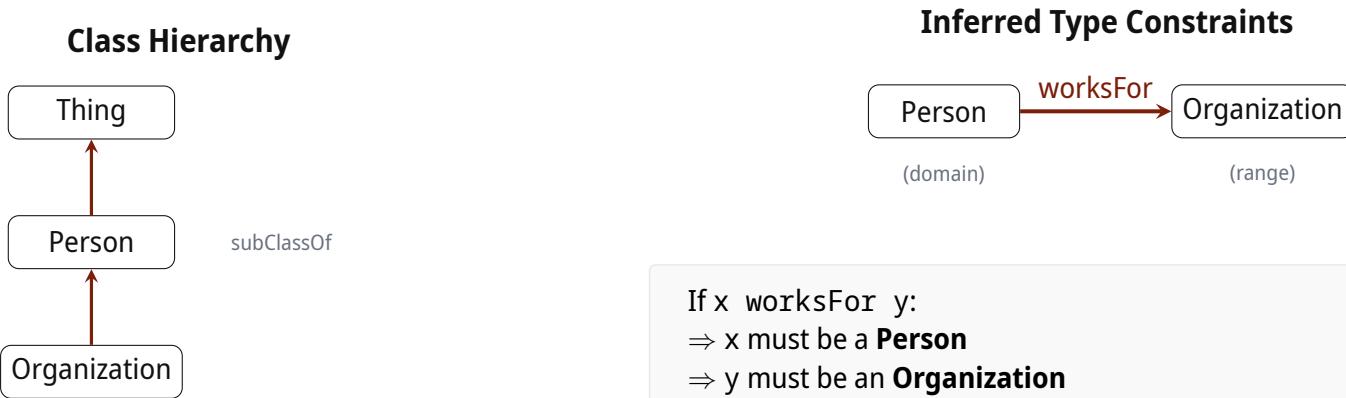
```
schema:Person a rdfs:Class ;  
rdfs:label "Person" ;  
rdfs:subClassOf schema:Thing .
```

**Relation** (Property)

```
schema:worksFor a rdf:Property ;  
domainIncludes schema:Person ;  
rangeIncludes schema:Organization .
```

*Source: Schema.org ontology*

# Hierarchy & Reasoning

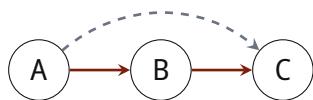


*Ontologies enable automatic type inference via first-order logic*

# Axioms & Inference

Axioms map to first-order logic, enabling inference of **implicit facts** from explicit data

## Transitivity



partOf

A partOf B, B partOf C

$\Rightarrow$  A partOf C

## Symmetry



knows

A knows B

$\Rightarrow$  B knows A

## Equivalence



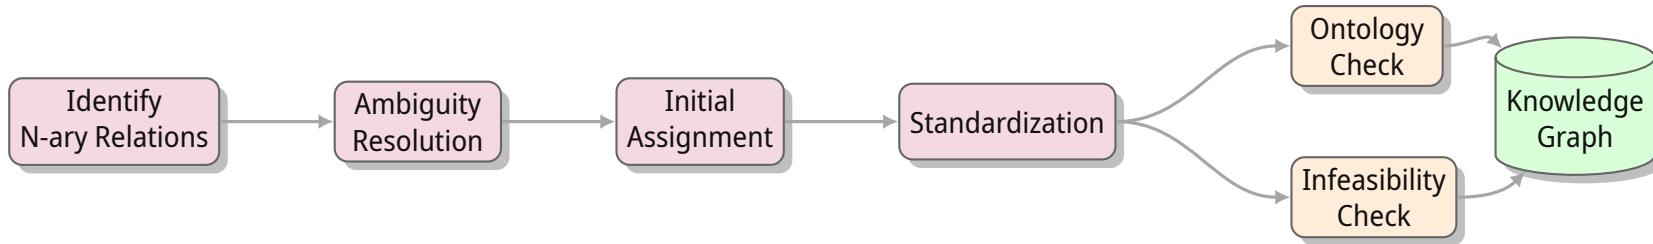
sameAs

A sameAs B

$\Rightarrow$  A  $\equiv$  B

**Example:** Given John worksFor Acme and worksFor has domain Person  
 $\Rightarrow$  Infer: John is a Person (implicit fact)

# KG Ingestion Pipeline



# N-ary Event Identification

- > LLM identifies **complex events** with 3+ participants
- > Events that can't be reduced to simple binary relations

## Manifest Storage

```
{  
  "short_id": "mtg_rome_01",  
  "description": "Meeting  
  between Victor & Maria  
  in Rome..."  
}
```

## Event VDB

Semantic search for existing events  
Prevents duplicate event creation  
Returns top-k similar events

*Events are reused across chunks via VDB similarity matching*

# Ambiguity Resolution & Initial Assignment



- > **Manifest Mandate:** LLM must use exact short\_id for identified events
- > **Initial Assignment:** LLM infers types (Person, Organization, Event...)

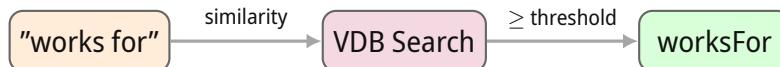
# Ontology & Standardization

## Types VDB

rdfs:Class entries  
Person, Organization,  
Location, Event...  
+ label, superclass

## Relations VDB

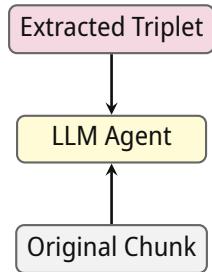
rdf:Property entries  
worksFor, locatedIn,  
memberOf, knows...  
+ domain, range



*Raw terms kept if similarity too low (no hallucination)*

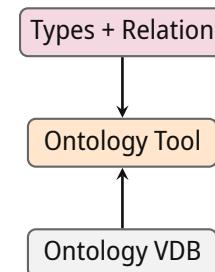
# Verification & Validation

## Chunk Verification



Supported? ✓ Keep  
Fabricated? ✗ Drop

## Ontology Feasibility

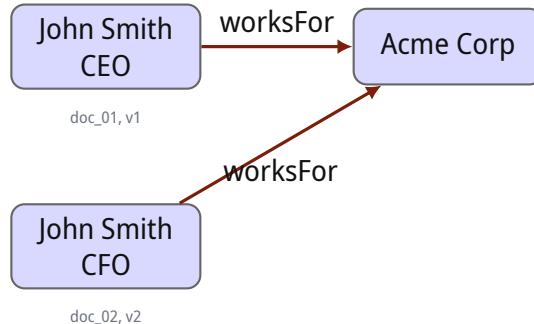


Valid combo? ✓ Keep  
Infeasible? ✗ Drop

*Double-check: both agent and tool must approve before KG insertion*

# Conflict Resolution & Versioning

> **Problem:** Same entity, different information across sources



**Solution:** New node with different metadata

Tracks: document\_id, chunk\_id, page\_no, ingested\_at

Preserves conflicting information with full provenance

# Beyond Storage

## KURSE: An Intelligent Data Understanding & Storage System

- > Highly adaptable architecture for diverse domains
- > Neuro-symbolic reasoning enables:
  - > Constraint Satisfaction
  - > Graph Traversal
  - > Logical Inference

*Let's see some applications...*

# Application: Legal Contradiction Detection

- > **The Idea:** Ingest case files (emails, receipts, testimonies) to verify if a suspect's alibi is physically possible
- > **Why KURSE Fits:**
  - > Treats this as a **Constraint Satisfaction Problem**
  - > Maps entities to Location(Time) states
  - > Applies rule: Entity cannot be at A and B at Time T
  - > Catches contradictions RAG misses (e.g., "at home" vs receipt at bar)
- > **Required Changes:** Change the Agentic Search's Query Answering agent to find conflicts between given info and retrieved info, and break the query into sub-queries based on the events claimed to happen in the input

# Application: Market Risk Identification

- > **The Idea:** Predict how localized events (e.g., Thailand flood) impact portfolio assets via supply chain effects
- > **Why KURSE Fits:**
  - > Treats this as a **Graph Traversal Problem**
  - > Builds supply chain: Flood → Factory → Supplier → Apple
  - > Propagates risk through transitive reasoning
  - > Finds 2nd/3rd order effects that keyword search misses
- > **Required Changes:** Add multi-hop traversal with risk propagation. Again, a minor change to the part of the Agentic Search that creates the sub-queries.



# BeeKurse

# What is BeeKurse?

## Conversational E-commerce Platform

Excelling in product recommendations

### > OCR + Reasoning for Small Vendors:

- > Update inventory by writing on paper daily
- > No need for detailed catalogued updates

### > VDB/KG Structure Enables:

- > Smart product recommendations
- > Product comparisons

Let's discuss the changes we made to  
KURSE

# Adapting KURSE for Commerce

- > **Dropped** rigid ontology system
- > **Replaced with** VDB seeded with relation types
- > **Why?** Commerce needs:
  - > Creative, particular relations
  - > Less rigid reasoning
  - > Flexibility over formality

# Search Algorithm: Query Classification



## 5 Query Types:

- > **SEARCH** — Product discovery
- > **DETAIL** — Product info Q&A
- > **CHAT** — Conversational
- > **CART\_ACTION** — Add/remove items
- > **CART\_VIEW** — View cart/wishlist

# Query Type: SEARCH — Parsing

## What it is:

- > User wants to **find new products**
- > Most complex query type

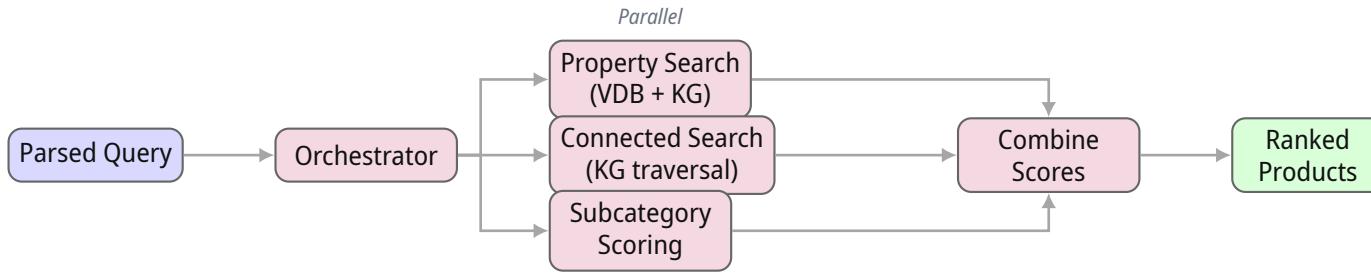
## Example:

"Red cotton shirt under \$30"

## Parsed Output:

```
{  
  "query_type": "SEARCH",  
  "category": "clothing",  
  "subcategory": "shirt",  
  "properties": [  
    ["red", 1.5, "HAS_COLOUR"],  
    ["cotton", 1.2, "HAS_MATERIAL"]  
  ],  
  "literals": [  
    ["price", "<", 30.0, 0.1]  
  ]  
}
```

# Query Type: SEARCH — After Parsing



- > **Property Search:** VDB similarity + KG relation matching
- > **Connected Search:** Find related products via KG edges
- > **Combine:** Weighted sum of all scores → rank → return top-K

# Query Type: DETAIL — Parsing

## What it is:

- > User asks about **specific product(s)**
- > Can reference by ID or short code
- > Supports comparisons

## Example:

"Compare the materials of p-123 and p-456"

## Parsed Output:

```
{  
  "query_type": "DETAIL",  
  "product_ids": [  
    "p-123",  
    "p-456"  
  ],  
  "original_query":  
    "Compare the materials"  
}
```

# Query Type: DETAIL — After Parsing



- > **Resolve:** Short codes (e.g., "44QM") → full product IDs
- > **SQL Fetch:** Retrieve complete product metadata
- > **LLM:** Product data as JSON context → answer user's question

# Query Type: CHAT

## What it is:

- > Greetings, general questions
- > No product search needed

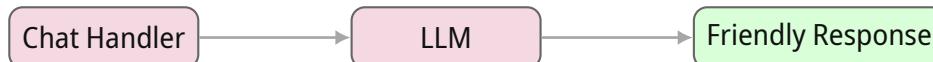
## Example:

"Hello!", "How does this work?"

## Parsed Output:

```
{  
  "query_type": "CHAT",  
  "message": "Hello!"  
}
```

## After Parsing:



No database queries — direct LLM response

# Query Type: CART\_ACTION

## What it is:

- > Add/remove items from cart
- > Manage wishlist

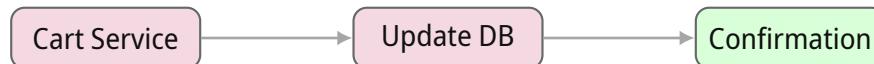
## Example:

"Add p-123 to my cart"

## Parsed Output:

```
{  
  "query_type": "CART_ACTION",  
  "action": "add",  
  "target": "cart",  
  "product_ids": ["p-123"]  
}
```

## After Parsing:



# Query Type: CART\_VIEW

## What it is:

- > Display cart or wishlist contents
- > Show product details

## Example:

"What's in my cart?"

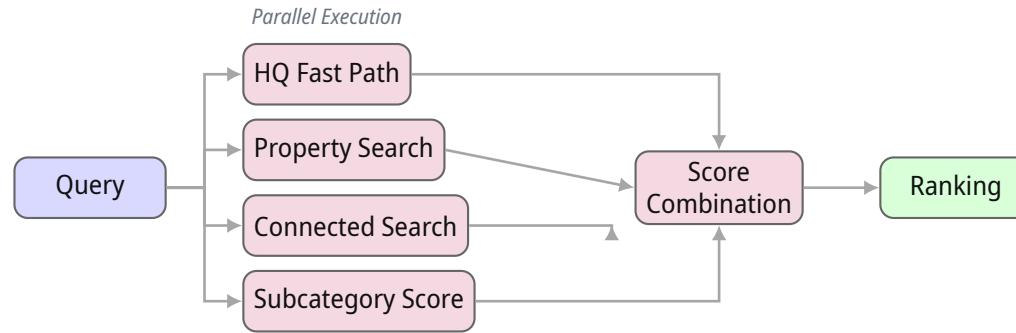
## Parsed Output:

```
{  
  "query_type": "CART_VIEW",  
  "target": "cart"  
}
```

## After Parsing:



# Search Orchestration Pipeline



Four parallel scoring paths, each contributing to final product ranking

- > **HQ:** Fast path for repeat purchases
- > **Property Search:** Semantic + KG matching
- > **Connected Search:** Recommendations via graph
- > **Subcategory:** Type-matching bonus

# HQ Fast Path (Hurry Query)

## What it is:

- > Direct SQL lookup for exact product
- > Bypasses all other search paths
- > Returns immediately if found

## When applicable:

- > `is_hq = True` (parser flag)
- > `prev_productid` exists
- > E.g., "my usual order"

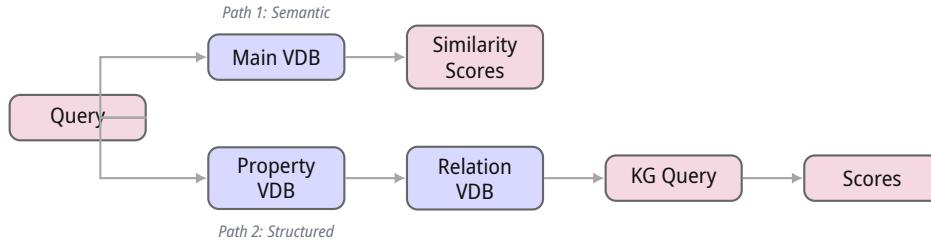


## Why we use it:

- > Instant response for repeats
- > Skips expensive VDB/KG queries
- > Best user experience

# Property Search (RQ — Relevance Query)

**What it is:** Main search path using VDB + KG



**When applicable:**

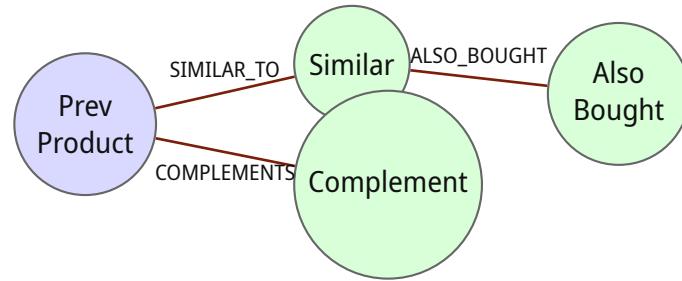
- > Always runs for SEARCH queries
- > (Unless HQ succeeds first)

**Why we use it:**

- > Semantic: "crimson" ≈ "red"
- > KG: structured property matching

# Connected Search (SQ — Similarity Query)

**What it is:** KG graph traversal to find related products



**When applicable:**

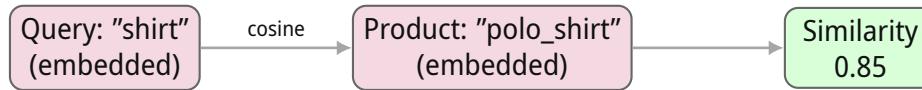
- > prev\_productid provided
- > prev\_storeid provided
- > "Show me similar to this"

**Why we use it:**

- > Recommendations from patterns
- > Store loyalty bonus (+0.2)
- > Connected bonus (+0.5)

# Subcategory Scoring

**What it is:** Cosine similarity between query and product subcategories



**When applicable:**

- Always runs
- Adds bonus if similarity > 0.5

**Why we use it:**

- "shirt" ranks shirts > pants
- Even if pants match "red cotton"
- Up to +0.4 bonus

# Property Weighting

## Weight Scale:

0.5 = Nice-to-have

1.0 = Normal importance

1.2-1.5 = Important

2.0 = Critical / must-have

## Formula:

score = similarity  $\times$  weight

## Example: "red cotton shirt"

### Properties:

"red" weight=1.5 (important)

"cotton" weight=1.0 (normal)

### Product A (red shirt):

red\_sim=0.95  $\rightarrow$   $0.95 \times 1.5 = 1.43$

cotton\_sim=0.80  $\rightarrow$   $0.80 \times 1.0 = 0.80$

### Product B (maroon shirt):

red\_sim=0.70  $\rightarrow$   $0.70 \times 1.5 = 1.05$

cotton\_sim=0.85  $\rightarrow$   $0.85 \times 1.0 = 0.85$

# Score Combination & Final Ranking

## Final Score Formula:

```
final_score = Σ(property_scores) + connected_bonus + subcategory_bonus -  
literal_penalties
```

### Bonuses

Connected: **+0.5**  
Same store: **+0.2**  
Subcategory: **up to +0.4**

### Literal Penalties

Buffer zone: 10%  
\$31 on \$30 limit:  
small penalty (-0.05)  
\$35 → filtered out

### Superlatives

"cheapest", "best rated"  
70% relevance  
30% literal value

**Example:** Product A: props=2.23 + connected=0.5 + subcat=0.34 - penalty=0 = **3.07**

# Thank You

## Contributions

### > Preprocessing, OCR & Parsing

Abhinav Goyal

### > Dockerization

Kunjan Manoj

### > Knowledge Graph Ingestion

Varshith Kada, Siripuru Abhiram

### > VDB Ingestion

Siripuru Abhiram, Varshith Kada

### > SQL Ingestion

Vishnu Teja

### > Search

Varshith Kada

### > Strontium Agent

Varshith Kada, Himesh

### > Frontend for Vendors

Himesh

### > Backend

Vishnu Teja

### > Synthetic Data Preparation

Bhuvan